

Lectures 5-7 Inf2C - Computer Systems: Intro to C

Boris Grot

School of Informatics
University of Edinburgh




Previous lectures

- MIPS
 - Arithmetic and memory
 - Control flow: branches and jumps
 - Function calls and the stack

Lectures 5-7: Intro to C

- Motivation:
 - C is both a high and a low-level language
 - Very useful for systems programming
 - Fast! (next slide)
- This intro assumes knowledge of Java
 - Focus is on differences
 - Most of the syntax is the same
 - Most statements, expressions are the same

Performance: C vs. the rest



Source: <http://attractivechaos.github.io/plb/>

Outline

- A simple program; how to compile and run
- Major differences with Java
- Data types and composite data structures
- Arrays and strings
- Pointers
- Other issues
 - Memory regions
 - C Preprocessor
 - Portability



The hello world program

```
#include<stdio.h>
```

```
int main(void)
{    // This is a comment
    printf("Hello world!\n");
    return 0;
}
```

Linux/DICE shell commands

Compile: `gcc hello.c`

Run: `./a.out`



Major differences with Java

- C is not object oriented
 - C programs are collections of **functions**, like Java methods, but not class-based.
 - No inheritance, subtyping, dynamic dispatch in C
- C is not interpreted
 - A C program is **compiled** into an executable machine code program, which runs directly on the processor
 - Java programs are compiled into a **byte code**, which is read and executed by the Java interpreter (which is just another program)

C is less “safe”

- Run-time errors are not ‘caught’ in C
 - The Java interpreter catches these errors before they are executed by the processor
 - Example: array out-of-bounds exception
 - C run-time errors happen for real and the program crashes
- The C compiler trusts the programmer!
 - Many mistakes go un-noticed, causing run-time errors and leaving systems vulnerable to security exploits

Memory management is different

- In Java
 - All objects dynamically allocated
 - Unusable objects recycled automatically by garbage collection
- In C
 - No objects, only data structures
 - Some data structures statically allocated, others dynamically
 - Dynamically-allocated storage must be reclaimed (or freed) once the data structures there are no longer needed.
 - Major source of error, particularly when the programmer forgets to free the memory, resulting in memory leaks.

C has pointers ...

- Pointers are special variables that reference (or point to) another variable
 - Similar to Java references
- We have already seen pointers in assembly:
`lw $t1, 0($s2)`
 - `$s2` is a pointer
 - C pointers are the same thing! (more later)

Built-in data types

- The usual basic data types are there:

char	8 bits
short	16
int	16, 32, 64 (same as machine word size)
long	32, 64
float	32
double	64
- Data type sizes are machine dependent
 - Unlike Java where an int is always 32 bits
- Normally signed, unsigned available too
- No boolean type exists
 - for any number (int, char,...): 0 false, other true

Composite data structures - struct


- Structures are like objects, but their types have no methods, unlike classes:

```
struct point {  
    int x, y;  
    // can include other structs  
} p1;  
struct point p2;
```

- Components accessed using “.” operator
`p1.x = 2;`

In memory: structures

```
struct point {  
    int x;  
    int y;  
} p1;
```



`sizeof(point) = 8`

What does `p1.y` translate into in MIPS?

```
addi $t0, $s0, 4 // $s0 points to the starting addr of p1  
lw   $t4, $t0    // load p1.y into $t4
```

User-defined types

- Define names for new or built-in types
`typedef <type> <name>;`

- Example:

```
typedef unsigned char byte;
```

```
typedef struct {
```

```
    inx x;
```

```
    int y;
```

```
} point;
```

```
...
```

```
point p1, p2;
```

Arrays

- Syntax of C arrays similar to Java
- As in Java, C arrays have fixed size
- Example declarations of array:

```
int m[] = {5, 8, 10}; // size fixed to 3
int n[2][10]; // two-dimensional array
                // with 2 rows and 10 cols
point p[4]; // array of 4 structs
```
- C arrays have no knowledge of their length
 - No checking that indexes are within bounds
- In C, close relationship between arrays and pointers
 - Pointers commonly used to pass arrays between functions

Strings

- C strings are simply arrays of type `char`
 - Encoded in 8 bits using ASCII
- They end with `'\0'`, the **null** character
 - `char s[10]; // up to 9 characters long`
- String initialisation
 - `char s[10] = "string"; // '\0' implied`
 - `char s1[] = "string, too"; // length= ?`
- Usual C rule for arrays apply:
 - Cannot store more chars than reserved at declaration
 - But bounds are not checked!

Strings – common operations

- Assignment: `strcpy(s, "string");`
- Length: `strlen(s)`
- To get the 6th character: `s[5]`
 - First char at position 0, as in Java arrays
- Comparison, `strcmp(s1, s2)` returns:
 - 0 when equal
 - Negative number when lexicographically $s1 < s2$
 - Positive when $s1 > s2$
- Must `#include<string.h>` to call the functions
 - Type: `man string` to see what's available




Pointers

- We have seen pointers in assembly:

```
lw $t1, 0($s2)
```

- \$s2 points to the location in memory where the “real” data is kept
- \$s2 is a register, but there’s nothing stopping us to have pointers stored in memory like “normal” variables



C pointers

- A C pointer is a variable that holds the address of a piece of data
- Declaration:
`int *p; // p is a pointer to an int`
 - The compiler must know what data type the pointer points to
- Basic pointer usage:
`p = &i; // p points to i now`
`*p = 5; // *p is another name for i`
- `&` - *address of* operator. `*` *dereference* operator

Pointers as function arguments

- In Java
 - an argument with primitive type is passed by value (function gets copy of value)
 - an argument with class type is passed by reference (function gets reference to value)
- In C
 - All arguments passed by value
 - To get effect of 'pass by reference', use an argument with a pointer type

Example – the swap function

```
void swap_wrong(int a, int b) {  
    int t=a;  
    a=b; b=t;  
}
```

swap_wrong swaps the local variables a, b which are unknown outside of the function

```
void swap(int *a, int *b) {  
    int t=*a;  
    *a=*b; *b=t;  
}
```

Function call: swap(&x, &y);

Pointer arithmetic and arrays

C allows arithmetic on pointers:

```
int a[10];
```

```
int *p;
```

```
p = a; // p points to a[0]. Same as p = &a[0]
```

p+1 points to a[1]

– Note that ~~&a[1] = &a[0]+1~~

– The compiler multiplies +1 with the data type size

In general: $p+i$ points to $a[i]$, $*(p+i)$ is $a[i]$

Also valid: $*(a+i)$ and $p[i]$

– but cannot change what a points to. It's not a variable

More pointer arithmetic

Common expressions:

- `*p++` use value pointed by `p`, make `p` point to next element
- `*++p` as above, but increment `p` first
- `(*p)++` increment value pointed by `p`, `p` is unchanged
- Special value `NULL` used to show that a pointer is not pointing to anything
 - `NULL` is typically 0, so statements like `if (!p)` are common
- Dereferencing a `NULL` pointer is a very common cause of C program crashes

Example – pointer arithmetic

Return the length of a string:

```
int strlen(char *s)
{
    char *p=s;
    while (*s++ != '\0');
    return s-p-1;
}
```

- Argument/variable *s* is local, so we can change it
- Pointer increment, dereference and comparison all in one! No statement in the loop body
- Note pointer subtraction at return statement

More fun with strings & pointers

```
char s1[10] = "Bob";
```

```
char s2[10] = "Bob";
```

```
if (s1 == "Bob")
```

```
    // do x
```

```
else if (s1 == s2)
```

```
    // do y
```

```
else
```

```
    // do z
```

Which statement (x, y, or z) is executed?

Dynamic memory allocation

- Pointers are not much use with **statically allocated** data
- Library function `malloc` allocates a chunk of memory at run time and returns the address

```
int *p;
if ((p = malloc(n*sizeof(int))) == NULL) {
    // Error
}

...
free(p); // release the allocated memory
```

Pointers to pointers

- Consider an array of strings:
`char *strTable[10];`
- The strings are **dynamically allocated** \Rightarrow any size
- But the table size is fixed to 10 strings
- What if we don't know the number of strings ahead of time?
 - Need to be able to provision array size on demand
 - That is, need to dynamically allocate the storage for the array of strings

`char **strTable;`



Pointers to pointers - details

Space must be allocated both for the table and the strings themselves

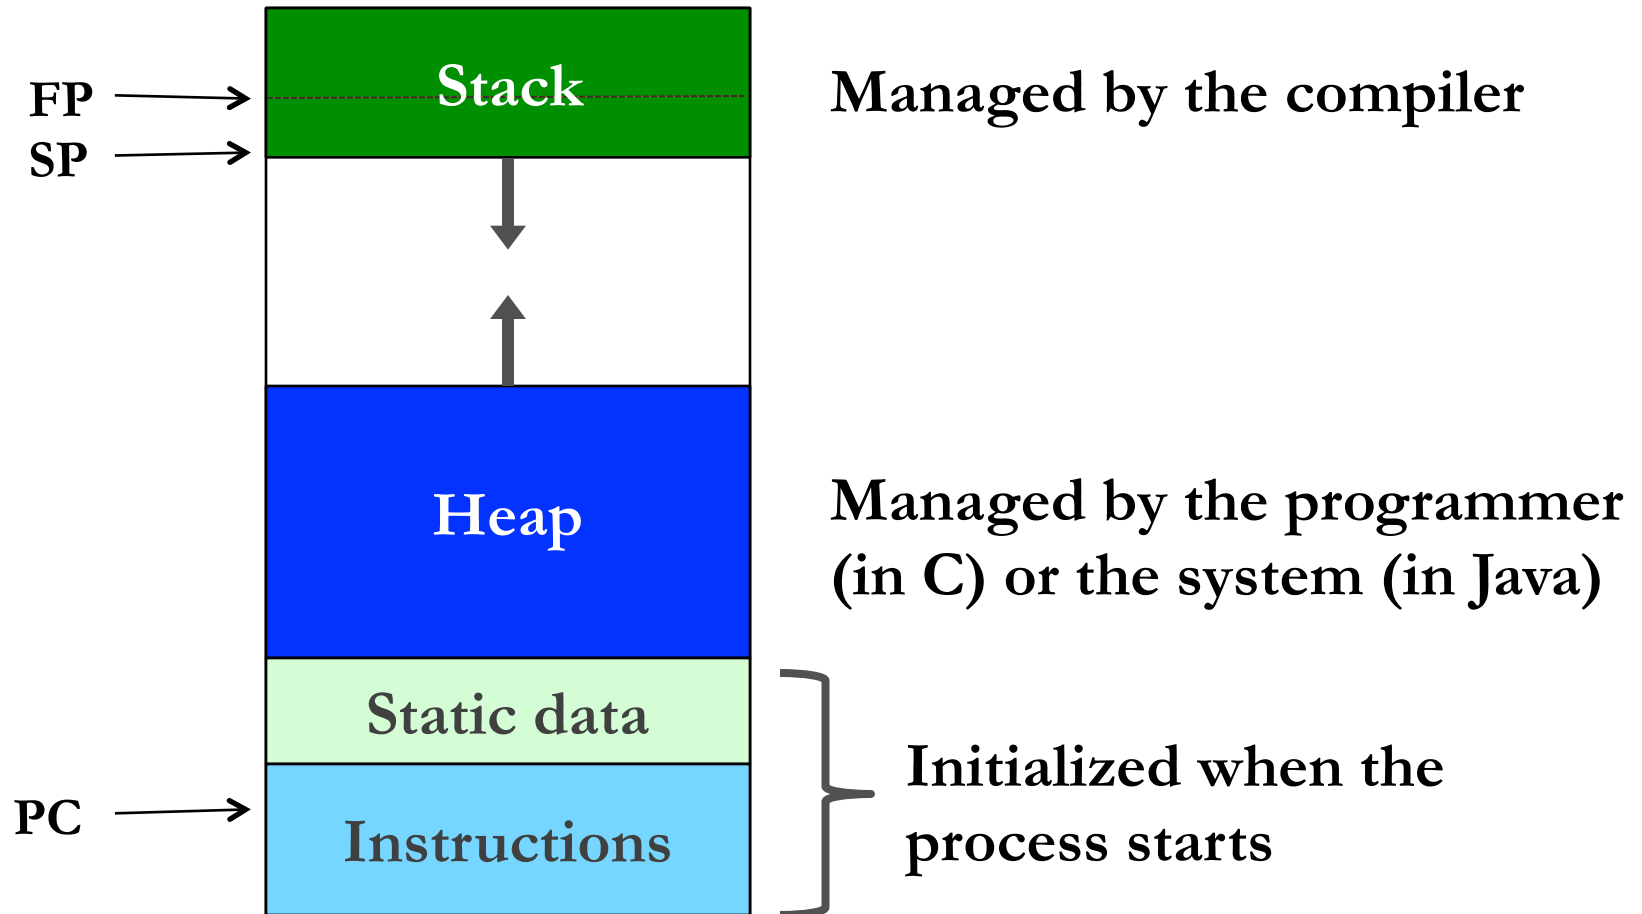
- Pointer to pointer!

```
1 char **strTable;
2 strTable = malloc(n*sizeof(char *));
3 for (i=0; i < n; i++) {
4     // s gets a string of length l
5     *(strTable+i) = malloc(l*sizeof(char));
6     strcpy(strTable[i], s);
7 }
8 // strTable[i][j] == (*(strTable+i)+j)
```

Memory regions and management

- Memory areas
 - *Heap*: dynamically allocated storage
 - *Stack*: for function/method local variables
 - *Static*: for data living program lifetime
- In Java
 - All objects on heap
 - Unusable objects on heap recycled automatically by garbage collection
- In C
 - Data structures in all 3 areas
 - Programs must explicitly free-up heap storage that is no longer needed

Memory regions in detail



Categories of variables in C

- Global variables (statically allocated)
 - Defined outside of functions
 - Have *lifetime* of program and *scope* to file end
 - **extern** declarations extend scope before definition and to other files
 - Declare **static** to hide from other files
- Local (*automatic*) variables (allocated on stack)
 - Defined inside a function
 - Not available outside function
 - Distinct storage for each function invocation
 - Declare **static** for same storage for all invocations

Compilation units

- Programs are divided into *compilation units*
 - Provide degree of modularity
 - Each commonly has main file (.c) for source code
 - *Header* files (.h) **declare** public interfaces of units
- Each compiled separately to relocatable object code
 - Allows creation of object-code libraries
- A *linker* assembles these into an *executable*, resolving references between units
- A *loader* sets up the executable program in memory and initialises data areas, prior to program being run
 - Loader also computes addresses for Jump instructions

Declaration vs Definition

- Declaration: inform the compiler of the existence of a variable or function

```
void swap(int *a, int *b);    // in .h file
```

- Definition: provide function body; allocate memory for globals

```
void swap(int *a, int *b) {    // in .c file
    int temp = *a;
    *b = a;
    *a = temp;
}
```

Compilation units example

A.h:

```
int array_len;           // global
extern int MAX_SIZE;     // global, defined elsewhere

// function declarations
void swap(int *a, int *b);
```

A.c:

```
#include "A.h"

// function definition
void swap(int *a, int *b) {
    int temp = *a;
    *a = *b;
    *b = temp;
}
```

```
main.c:
#include <stdio.h>
#include "A.h"

int main(void) {
    int a = 5;
    int b = 15;
    swap(a, b);
}
```

Error?



The C pre-processor: `cpp`

- Includes – imports header files
`#include <stdio.h>`
`#include "A.h"`
 - Text substitution, e.g. define constants
`#define NAME value`
 - Macros (inline functions)
`#define MAX(X,Y) (X>Y ? X : Y)`
 - Conditional compilation
`#ifdef DEBUG`
`Printf("Debugging message");`
`#endif`
- > `gcc -DDEBUG ...`



That's all folks

- Not all C features have been covered, but this introduction should be enough to get you started
- Useful things to learn on your own:
 - Standard input/output: `printf`, `scanf`, `getc`, ...
 - File handling: `fopen`, `fscanf`, `fprintf`, ...
- Look over past exam papers for simple C programming exercises

Coursework 1

- Assigned “now”, due in 2 weeks
 - **Deadline: Tue, 27 Oct, 16:00h**
- Task A: split a character string into words
 - Given: a C implementation
 - Your job: convert it to MIPS
- Task B: find single-word palindromes in a string
 - Given: C and MIPS implementations of Task A
 - Your job: write C and MIPS code for Task B



Coursework 1 (con'd)

- Task A example:

input: The first INF2C-CS coursework

output:

The

first

INF2C

CS

coursework

- Task B example:

input: I got my Honda Civic in 2002.

output:

Civic

2002



A (friendly) note on plagiarism

- **Don't do it!!!!**
- We use MOSS to electronically cross-check all submissions
 - MOSS is unaffected by variable renaming, code reshuffling, etc.
- Two plagiarism instances (4 students in total) were detected and prosecuted last year.
 - Remember: if you're sharing your code, you're just as guilty as the person taking it.

