

Inf2C - Computer Systems

Lecture 14

Virtual Memory

Boris Grot

School of Informatics
University of Edinburgh



Previous lecture: Memory hierarchy

- Main idea: exploit locality in memory references to create the illusion of a fast & large memory
 - Temporal vs spatial locality
- Memory hierarchy levels: registers, cache (≥ 1 levels), main memory, disk
- Cache: hardware-managed storage
 - Exploits temporal & spatial locality
 - Fully-associative vs direct mapped

Lecture 14: Virtual memory

- Motivation
- Overview
- Address translation
- Page replacement
- Fast translation – TLB

Motivation

Virtual memory addresses two main problems:

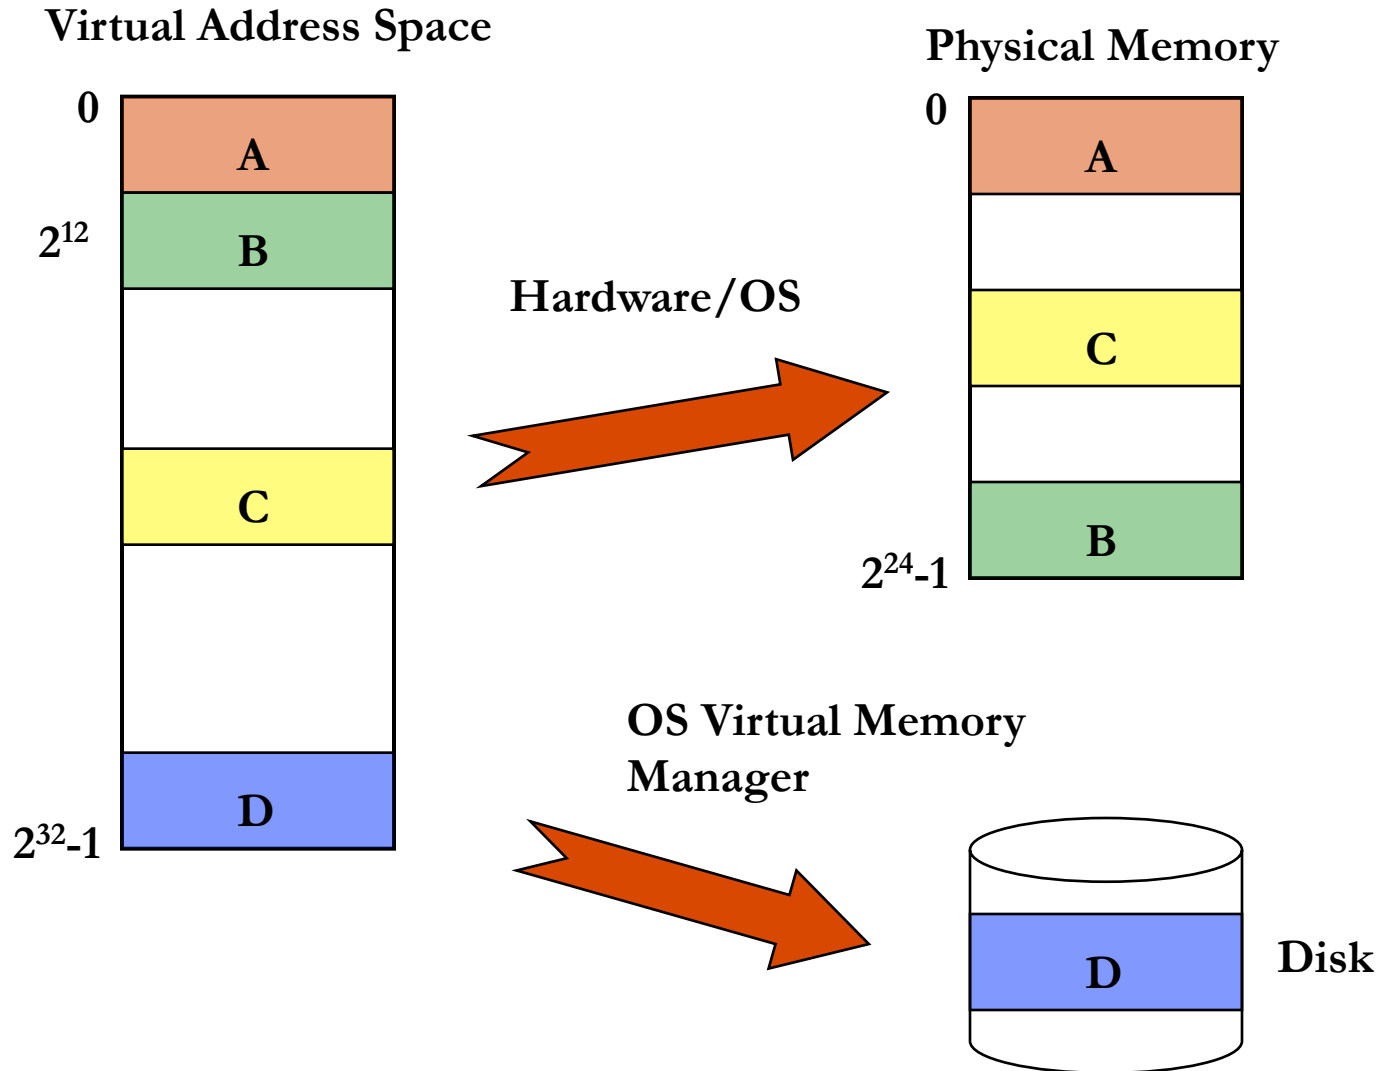
- 1) Capacity: how do we remove burden of programmers dealing with limited main memory?
 - Want to allow for the physical memory to be smaller than the program's **address space** (e.g., 32 bits → 4GB)
 - Want to allow multiple programs to share the limited physical memory with no human intervention
- 2) Safety: how do we allow for safe and efficient sharing of memory among multiple programs?
 - Want to prevent user programs from accessing the memory used by the OS
 - Want strict control of access by each user program to memory of other user programs

Virtual Memory

- Basic idea: each program thinks it owns the entire memory → the **virtual address space**
 - PC and load/store addresses are **virtual addresses**
- Actual main memory: **physical address space**
 - Virtual addresses are **translated** on-the-fly to physical addresses
 - Parts of virtual address space not recently used are stored on disk
- Dynamic address translation is done by combination of hardware and the OS



Address translation for 1 process



Physical memory as cache for VM

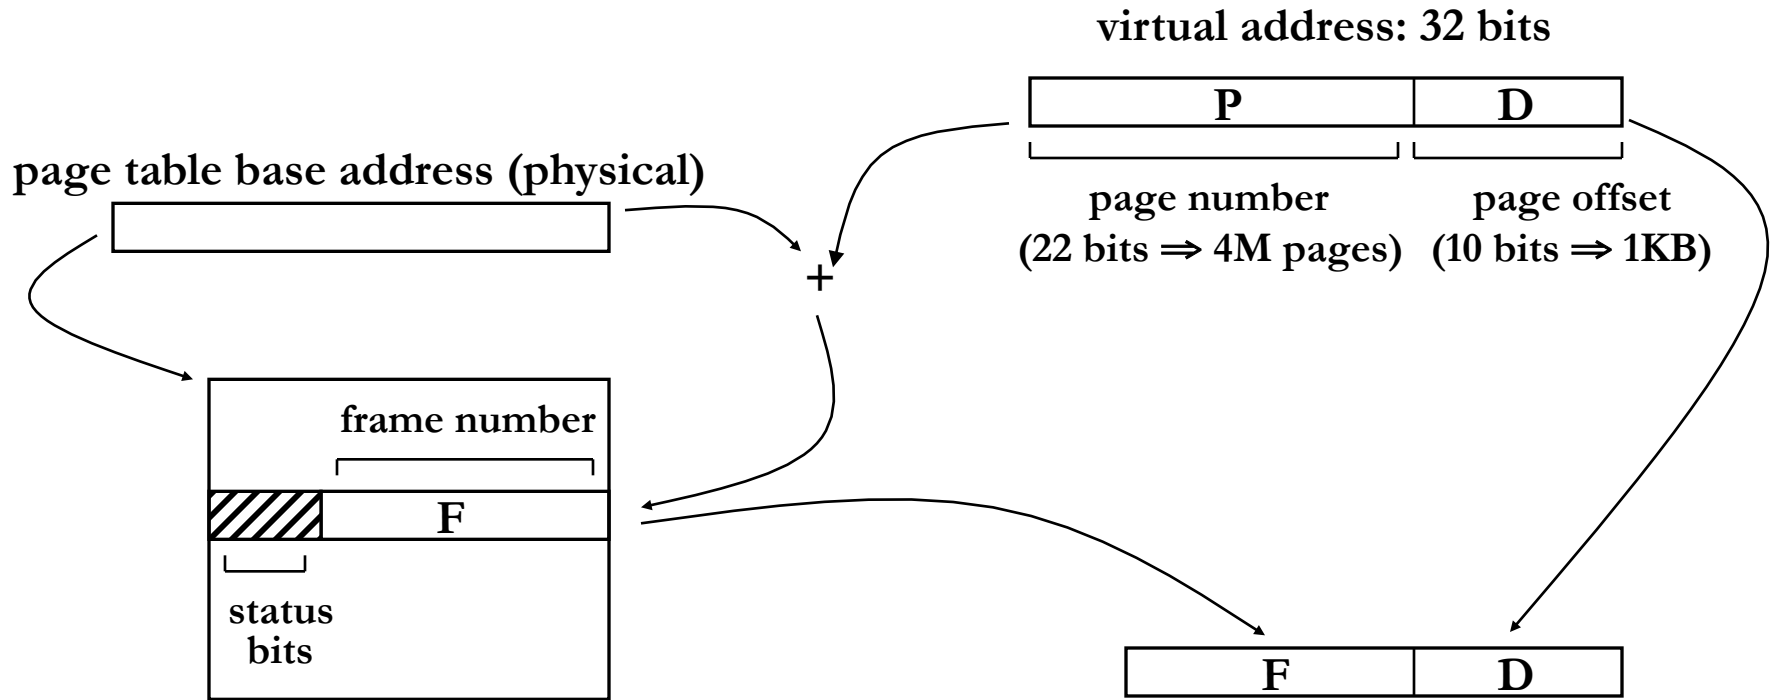
- Virtual memory space can be larger than physical memory
 - Programmer always sees the full address space (MIPS: 2^{32} bytes)
- Physical memory used as a cache for the virtual memory
 - Physical memory holds the currently used portions of a process' code and data areas (exploits locality!)
- Secondary storage (usually disk) “backs” the physical memory
 - OS reserves a portion of the disk for **swap space**
 - OS swaps portions of each process' code and data areas in and out of physical memory on demand (process called **paging**)
 - Swapping is transparent to the programmer



Paging

- A “cache line” or “block” of VM is called a **page**
 - Plain “**page**” or “**virtual page**” for virtual memory
 - “**Page frame**” or “**physical page**” for physical memory
- Typical sizes are 4-8 KB (MB or GB in servers)
 - Large enough for efficient disk use and to keep translation tables small
- Mapping is done through a per-process **page table**
 - Allows control of which pages each process can access
 - Different processes can use same virtual addresses

Dynamic Address Translation



page table:

- per process
- one entry per page (e.g. 4M entries)
- located in the system portion of main memory

physical address: 30 bits
(1GB of main memory)

Moving pages to/from memory

- Access to a non-allocated page causes a **page-fault** which invokes the OS through the interrupt mechanism
 - **R**(esidence) bit in page table status bits is zero
- Pages are allocated on demand
- Pages are replaced and swapped to disk when system runs out of free page frames
 - Aim to replace pages not recently used (principle of locality). **A**(ccess) bit for a page is set whenever page is accessed and is reset periodically
 - If any data in page has been modified, the page must be written back to disk: **M**(odified) bit in status bits is set

Providing Protection


- Each page table entry can have permission bits controlling
 - access allowed or not
 - read & write, read only or execute only access
- This enables per-process memory protection
 - E.g. can set up private and shared areas
- Important that only OS can change page tables
 - One of motivations for having kernel & user modes

Translation Lookaside Buffer

- Accessing the page table is costly in terms of latency
 - Two memory accesses per load and store (1 to get the page table entry + 1 to get the data)
- Fast address translation: **Translation Lookaside Buffer (TLB)** contained in the MMU
 - Small and fast table in hardware, located close to processor
 - Is a cache for page table: holds frame addr., not program data
 - Tag: virtual address. Entry: physical frame address
 - Can capture most translations due to principle of locality
 - When page not in TLB: access the page table, and save the translation entry in TLB

Translation Lookaside Buffer

Fully associative memory:



Virtual Memory: full picture

