

Inf2C - Computer Systems

Lecture 15

Exceptions and Processor Management

Boris Grot

School of Informatics

University of Edinburgh



Previous lecture: Virtual memory

- Solves two problems:
 - Capacity (physical memory is limited)
 - Safety (physical memory must be shared by multiple programs and the OS)
- Virtual vs physical address space
 - Each program “sees” a full 32-bit address space
 - Actual physical memory managed by the OS
- Address translation
 - Page table – all translations, but slow (in memory)
 - TLB – recent entries only, but fast (cache)



Exceptions – definition

- Exceptional events that interrupt normal program flow and require attention of the CPU
- External (“interrupts”)
 - Not caused by program execution
 - E.g. I/O interrupt (e.g., network packet arrived)
- Internal (“traps”)
 - Caused by program execution
 - E.g. illegal instruction, arithmetic overflow, TLB miss

Intentional exceptions

- Use exception mechanism to request some OS functions
 - e.g., I/O (e.g., print to screen), memory allocation
- User program uses **syscall** instruction
 - Cause register (\$v0) is set with a special value to identify the syscall exception
 - OS exception handler invoked when instruction executes
- Parameters are passed to the OS through agreed upon registers (usually \$a0, \$a1, ..)



Syscall example

The following will print the integer in register \$t0 to the screen.

```
li    $v0, 1      # service 1 is "print integer"  
add   $a0, $t0, $zero # load integer into $a0  
syscall
```

Exception mechanism

- Step 1: Save the address of current instruction
 - into a special register, the **exception program counter** (EPC)
 - Note: must return to the interrupted instruction (not PC+4)
- Step 2: Transfer control to the OS at a known address (i.e., exception handler PC)
- Step 3: Handle the exception
 - Deal with the cause of the exception
 - All registers must be preserved, similar to a procedure call
- Step 4: Return to user program execution
 - Handler restores user program's registers and jumps back using EPC
 - Relies on special instruction **eret**



Finding the exception handler

- Approach 1:
 - Jump to a predefined address (0x800000180)
 - Use the **Cause** register to then branch to the right handler (e.g., print int, read string, exit program)
 - Works well for syscall – cause register explicitly set
- Approach 2
 - Directly jump to a specific handler depending on the exception (**vectored interrupt**)
 - Eg:

Undefined opcode:	0xC000 0000
Overflow:	0xC000 0020
....:	0xC000 0040

Handling the exception

- Determine action required
 - By inspecting the Cause register or by virtue of being at the right handler (e.g., undefined opcode)
- If restartable:
 - Take corrective action, then use EPC to return to program
- Otherwise:
 - Terminate program and report error using EPC, cause, ...
- For a critical time while the interrupt is being handled, other interrupts should not happen
 - Otherwise the EPC, Cause will be overwritten
 - This is forced by masking interrupts → by setting the **exception level** (EXL) bit in the **status** register



Protecting system resources

- The OS must guarantee safe and orderly access to critical system resources
 - Hardware (processor, networking, I/O)
 - Program memory (including page tables)
- The OS is the ultimate arbiter of what's allowed
 - TLB miss → OK (but must access page table to service)
 - Write access to a read-only page → not OK (but must access page table to check)
 - Illegal opcode → not OK (kill the program)
- Exceptions are used to hand control over to the OS
 - Need a separate mechanism to limit capabilities of user programs



Kernel vs. User Mode Protection

- Exceptions (including system calls) are handled by the OS
 - CPU has two modes of operation: **user** and **kernel** (OS)
 - Current mode identified by a bit in a special status register
 - Exception mechanism is used to force the mode to change from user to kernel for execution of OS functions
- “Privileged” instructions only executed in kernel mode
 - E.g. accessing I/O devices, handling page tables accesses
- Kernel mode can only be entered through an exception
 - User programs cannot jump to OS instruction space

iret instruction sets mode back to previous mode



Advantages of Dual Mode architecture

- Guarantees that control is transferred to OS when user programs attempt to perform potentially dangerous tasks
- Allows OS to ensure that programs do not interfere with each other
 - e.g., that memory is divided appropriately
- Allows OS to ensure that programs do not have access to resources for which they do not have permission
 - e.g., files
- Ensures that user programs do not have indefinite control of the processor
 - could happen in old-school Windows or ancient DOS



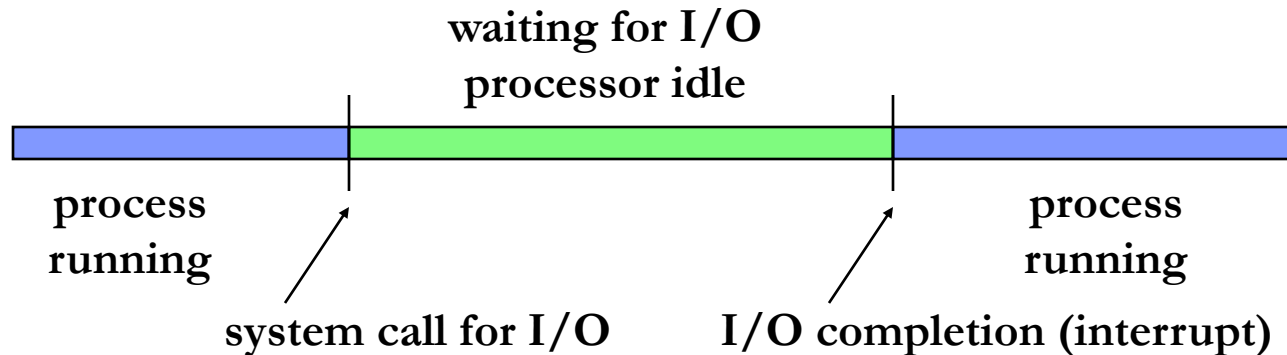
Managing the Processor

- Problem:
 - I/O takes too long → processor idle
 - User programs can crash or monopolize the CPU (either unintentionally or maliciously)
- Solution:
 - **Multiplex** or **time-share** the CPU and other resources among several user processes
 - Switch from one process to another when it performs I/O, or when its time allocation (time slice) expires

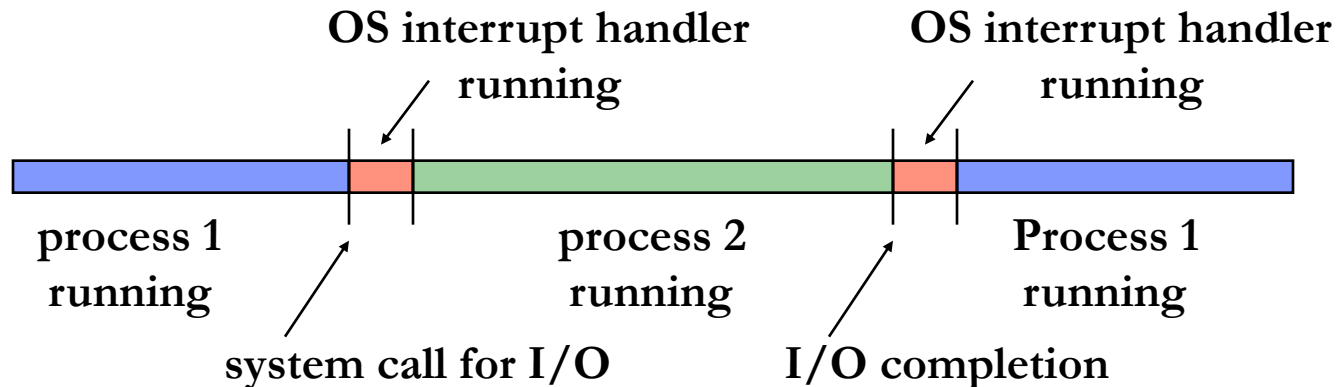
Process: “a program in execution” [Silberschatz, Galvin, Gagne]

Multi-tasking

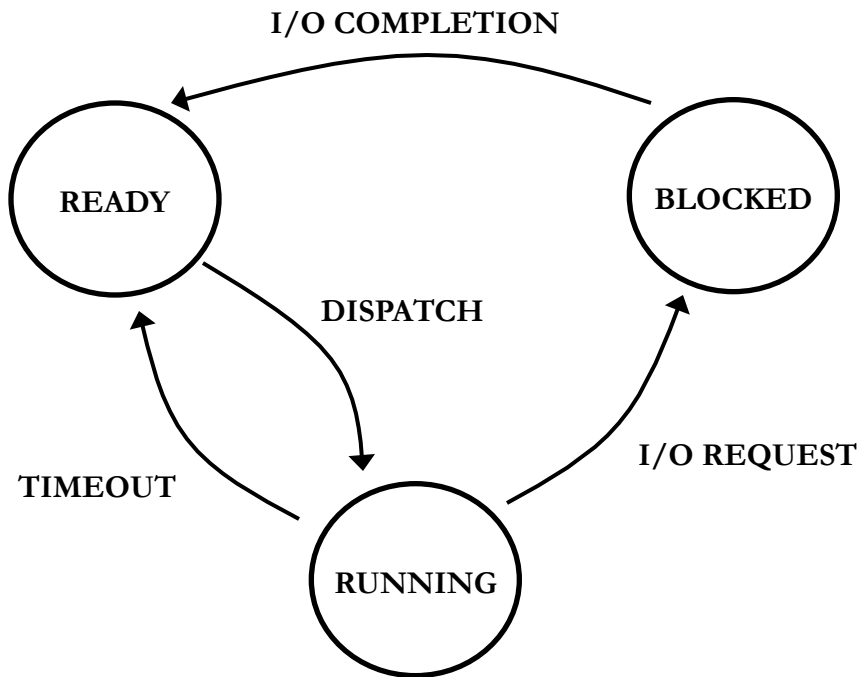
- Single-task system:



- Multi-tasking system:



Process States



States:

RUNNING: process is currently running in the CPU

READY: process is not running, but could run if brought into CPU

BLOCKED: process is not able to run because it is waiting for I/O to finish

Transitions:

I/O REQUEST: process initiates I/O

I/O COMPLETION: I/O finishes

DISPATCH: OS moves process into CPU and it starts executing

TIMEOUT: process's timeslice is over

Process States

- Step 1: process calls (or **traps into**) the OS, or interrupt occurs (e.g. because of timer)
- Step 2: OS's **dispatcher** performs **context-switch**:
 - Process's context is saved (registers, PC, etc) in **process control block** (PCB)
 - Dispatcher chooses new process to run
 - Processes' states are updated

PCB: OS data structure containing each process's information:

- Process id (PID)
- Process state (blocked, running, etc)
- Process priority
- Process permissions
- Etc



Suspending and Resuming Processes

■ Problem:

- Might not have enough physical memory for all processes
- Some processes have higher priority and must get more processor time (e.g., media playback)

■ Solution:

- Processes can be “swapped out” from memory to disk
- Such processes are moved into an “inactive” state
 - 2 new process states
- PCB of inactive processes are still kept in OS memory
- Inactive processes are resumed by “swapping in” the data from disk back to memory



Suspending and Resuming Processes

