

Adversarial Search


R&N: § 5.1-5.4

Michael Rovatsos

The logo for the School of Informatics, featuring a stylized 'i' with a red dot and the word 'informatics' in a bold, sans-serif font.
University of Edinburgh

19 January 2016

Informatics 2D




Outline

- Games
- Optimal decisions
- α - β pruning
- Imperfect, real-time decisions

Games vs. search problems

- We're (usually) interested in zero-sum games of perfect information
 - Deterministic, fully observable
 - Agents act alternately
 - Utilities at end of game are equal and opposite
- "Unpredictable" opponent → specifying a move for every possible opponent reply
- Time limits → unlikely to find goal, must approximate

Game tree (2-player, deterministic, turns)



- 2 players: MAX and MIN
- MAX moves first
- Tree built from MAX's POV

Utility of each terminal state from MAX's point of view.

Optimal Decisions

- Normal search: optimal decision is a sequence of actions leading to a goal state (i.e. a winning terminal state)
- Adversarial search:
 - MIN has a say in game
 - MAX needs to find a contingent strategy which specifies:
 - MAX's **move** in initial state then...
 - MAX's **moves** in states resulting from every response by MIN to the **move** then...
 - MAX's moves in states resulting from every response by MIN to all those **moves**, etc...


minimax value of a node = utility for MAX of being in corresponding state:

MINIMAX(s) =

$$\begin{cases} \text{UTILITY}(s) & \text{if } \text{TERMINAL-TEST}(s) \\ \max_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER}(s) = \text{MAX} \\ \min_{a \in \text{Actions}(s)} \text{MINIMAX}(\text{RESULT}(s,a)) & \text{if } \text{PLAYER}(s) = \text{MIN} \end{cases}$$

Minimax

- Perfect play for deterministic games
- Idea: choose move to position with highest minimax value = best achievable payoff against best play
- Example: 2-ply game:



Minimax algorithm

```
function MINIMAX-DECISION(state) returns an action  
  return  $\arg \max_{a \in \text{ACTIONS}(s)} \text{MIN-VALUE}(\text{RESULT}(\text{state}, a))$ 
```

```
function MAX-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a)))$   
  return v
```


```
function MIN-VALUE(state) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow \infty$   
  for each a in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a)))$   
  return v
```

Idea: Proceed all the way down to the leaves of the tree then
minimax values are backed up through tree


Properties of minimax

- **Complete?** Yes (if tree is finite)
- **Optimal?** Yes (against an optimal opponent)
- **Time complexity?** $O(b^m)$
- **Space complexity?** $O(bm)$ (depth-first exploration)
- For chess, $b \approx 35$, $m \approx 100$ for "reasonable" games
 - exact solution completely infeasible!
 - would like to eliminate (large) parts of game tree


α - β pruning example




α - β pruning example




α - β pruning example



α - β pruning example



α - β pruning example



Are minimax value of root and, hence, minimax decision **independent** of pruned leaves?

Let pruned leaves have values u and v ,

then $\text{MINIMAX}(\text{root}) = \max(\min(3, 12, 8), \min(2, u, v), \min(14, 5, 2))$

$= \max(3, \min(2, u, v), 2)$

$= \max(3, z, 2)$ where $z \leq 2$

$= 3$

→ Yes!

Properties of α - β

- Pruning does not affect final result (as we saw for example)
- **Good move ordering** improves effectiveness of pruning (How could previous tree be better?)
- With “perfect ordering”, time complexity = $O(b^{m/2})$
 - branching factor goes from b to \sqrt{b}
 - (alternative view) doubles depth of search compared to minimax
- A simple example of the value of reasoning about which computations are relevant (a form of **meta-reasoning**)

Why is it called α - β ?

- α is the value of the best (i.e., highest-value) choice found so far at any choice point along the path for MAX
- If v is worse than α , MAX will avoid it
 → prune that branch
- Define β similarly for MIN


MAX

MIN

..
..
..

MAX

MIN



The α - β algorithm

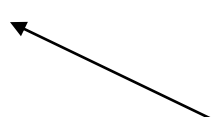
```
function ALPHA-BETA-SEARCH(state) returns an action  
   $v \leftarrow \text{MAX-VALUE}(\text{state}, -\infty, +\infty)$   
  return the action in ACTIONS(state) with value  $v$ 
```

```
function MAX-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MAX}(v, \text{MIN-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \geq \beta$  then return  $v$  ← Prune as this value is  
     $\alpha \leftarrow \text{MAX}(\alpha, v)$  worse for MIN and so  
  return  $v$  won't ever be chosen by  
  MIN!
```

α is value of the best i.e. highest-value choice found so far at any choice point along the path for MAX

β is value of the best i.e. lowest-value choice found so far at any choice point along the path for MIN

The α - β algorithm

```
function MIN-VALUE(state,  $\alpha$ ,  $\beta$ ) returns a utility value  
  if TERMINAL-TEST(state) then return UTILITY(state)  
   $v \leftarrow -\infty$   
  for each  $a$  in ACTIONS(state) do  
     $v \leftarrow \text{MIN}(v, \text{MAX-VALUE}(\text{RESULT}(s, a), \alpha, \beta))$   
    if  $v \leq \alpha$  then return  $v$    
     $\beta \leftarrow \text{MIN}(\beta, v)$   
  return  $v$ 
```

Prune as this value is worse for MAX and so won't ever be chosen by MAX!

Resource limits

Suppose we have 100 secs, explore 10^4 nodes/sec
→ 10^6 nodes per move

Standard approach:

- cutoff test:
 - e.g., depth limit (perhaps add quiescence search, which tries to search interesting positions to a greater depth than quiet ones)
- evaluation function
 - = estimated desirability of position

Evaluation functions

- For chess, typically linear weighted sum of features

$$\text{EVAL}(s) = w_1 f_1(s) + w_2 f_2(s) + \dots + w_n f_n(s)$$

where each w_i is a weight and each f_i is a feature of state s

- Example
 - queen = 1, king = 2, etc.
 - f_i = number of pieces of type i on board
 - w_i = value of the piece of type i

Cutting off search

Minimax Cutoff is identical to *MinimaxValue* except

1. *TERMINAL-TEST* is replaced by CUTOFF
2. *UTILITY* is replaced by EVAL

Does it work in practice?

$$b^m = 10^6, b=35 \rightarrow m=4$$

4-ply lookahead is a hopeless chess player!

- 4-ply \approx human novice
- 8-ply \approx typical PC, human master
- 12-ply \approx Deep Blue, Kasparov

Deterministic games in practice

- Checkers: Chinook ended 40-year-reign of human world champion Marion Tinsley in 1994. Used a precomputed endgame database defining perfect play for all positions involving 8 or fewer pieces on the board, a total of 444 billion positions.
- Chess: Deep Blue defeated human world champion Garry Kasparov in a six-game match in 1997. Deep Blue searches 200 million positions per second, uses very sophisticated evaluation, and undisclosed methods for extending some lines of search up to 40 ply.
- Othello: human champions refuse to compete against computers, who are too good.
- Go: human champions refuse to compete against computers, who are too bad. In Go, $b > 300$, so most programs use pattern knowledge bases to suggest plausible moves.

Summary

- Games are fun to work on!
- They illustrate several important points about AI
- perfection is unattainable → must approximate
- good idea to think about what to think about